AURIS Platform Architecture v1.0

Autonomous Unified Response Intelligence System

Abstract

Version: 1.0
Date: March 2026
Author: Suzanne Natalie Button, Director
Organisation: Intelligent Consulting BV, Brussels, Belgium
Status: Published

Table of Contents

Executive Summary 4
1. Introduction 5
1.1 Purpose 5
1.2 The SIEM Replacement Imperative 5
1.3 Scope 5
1.4 Target Audience 5
2. Architectural Overview 7
2.1 Six-Layer Architecture 7
2.2 Design Philosophy 7
3. Layer 1: STRATIX Telemetry Mesh 8
3.1 Overview 8
3.2 Architecture Components 8
3.3 Telemetry Sources 8
3.4 Cost Optimisation 9
3.5 Data Ownership and Control 10
4. Layer 2: Intelligence-First Detection 11
4.1 Overview 11
4.2 Behavioural Baseline Analysis 11
4.3 Continuous Machine Learning Detection 11
4.4 AI-Augmented Chain Correlation 12
4.5 Threat Intelligence Operationalisation 12
5. Layer 3: Autonomous Operations 13
5.1 Overview 13
5.2 Multi-Agent Orchestration Architecture 13
5.3 Governance Framework 13
5.4 The Aviation Autopilot Model 14
5.5 Performance Metrics 14
6. Layer 4: Unified Threat Operations (UTO) 15
6.1 Overview 15
6.2 Cross-Domain Correlation 15
6.3 OT/ICS as First-Class Participant 15
6.4 Cloud Security Posture Management (CSPM) Native 16
6.5 UTO Operational Workflow 16
7. Layer 5: Sovereign-by-Design Architecture 17

© Intelligent Consulting 2026

7.1 Overview 17
7.2 Sovereignty Dimensions 17
7.3 US CLOUD Act Risk 17
7.4 Sovereign Architecture Requirements 18
7.5 NIS2/DORA/GDPR/AI Act Alignment by Default 18
8. Layer 6: Continuous Adversarial Validation (CAV) 20
8.1 Overview 20
8.2 CAV Operational Model 20
8.3 Detection-as-Code 20
8.4 Operationalised Threat Intelligence 20
8.5 MITRE ATT&CK Coverage Metrics 21
9. Platform Implementation Models 22
9.1 AURIS Platform Agnosticism 22
9.2 Implementation Approaches 22
9.3 Vendors to Challenge 23
10. The Human Dimension: Redefining the Analyst Role 24
10.1 AURIS Does Not Eliminate Analysts 24
10.2 The AURIS SOC Operating Model 24
11. Economics and Business Case 26
11.1 Total Cost of Ownership Analysis 26
11.2 Business Case by Stakeholder 26
12. Implementation Roadmap 28
12.1 Four-Phase Transformation Programme 28
Phase 1: Foundation (Months 1–6) 28
Phase 2: Intelligence (Months 6–12) 28
Phase 3: Autonomy (Months 12–18) 29
Phase 4: Optimisation (Months 18–24) 29
13. Performance Metrics and SLAs 30
13.1 The New Metrics of Security Operations Excellence 30
14. Security and Resilience 31
14.1 Platform Security Architecture 31
14.2 Resilience and Business Continuity 31
15. Standards and Certifications 32
15.1 Compliance Framework Support 32
15.2 Certification Roadmap 32
16. Vendor Ecosystem and Partnerships 33
16.1 Strategic Technology Partners 33

© Intelligent Consulting 2026

17. Conclusion: The Strategic Imperative 34
References 35
Appendix A: AURIS Agent Governance Template 36
A.1 Autonomy Decision Matrix 36
A.2 Human-in-the-Loop Thresholds 36
Appendix B: AURIS Deployment Checklist 37
B.1 Pre-Deployment Requirements 37
B.2 Phase 1 Foundation Validation 37
B.3 Phase 2 Intelligence Activation 37
B.4 Phase 3 Autonomy Deployment 37
B.5 Phase 4 Optimisation Completion 38
Appendix C: Glossary 39

© Intelligent Consulting 2026

Executive Summary

AURIS (Autonomous Unified Response Intelligence System) is a sovereign-by-design, intelligence-first security operations platform that replaces the legacy SIEM stack with a coherent, continuously learning system built on six integrated layers[1]. Operating on the STRATIX Telemetry Mesh as its foundational schema, AURIS achieves 60–75% total cost of ownership reduction, 80% autonomous alert resolution, and native compliance with NIS2, DORA, GDPR, and the EU AI Act[1].

AURIS represents the complete architectural answer to SIEM obsolescence, delivering mean time to detect under 5 minutes, 70% reduction in false positives, and 80% MITRE ATT&CK coverage through continuous adversarial validation[1].

© Intelligent Consulting 2026

1. Introduction

1.1 Purpose

This technical specification defines the complete AURIS platform architecture, operational layers, autonomous agent framework, implementation patterns, and deployment models for modern security operations centres.

1.2 The SIEM Replacement Imperative

Legacy SIEM architectures are structurally obsolete[1]:

- €380k–€940k annual operational cost for in-house SIEM
- €5.6M–€9.4M total security stack cost at enterprise scale
- 40–60% false positive rates overwhelming analyst capacity
- Hours to days mean time to detect
- No native OT/ICS coverage
- Manual compliance reporting (1,200+ analyst hours per NIS2 incident)

AURIS eliminates these structural limitations through a fundamentally different architecture built around data, intelligence, and autonomous action rather than log storage and rule engines[1].

1.3 Scope

AURIS covers:

- Six-layer platform architecture
- STRATIX Telemetry Mesh implementation
- Intelligence-first detection methodology
- Autonomous operations agent framework
- Unified Threat Operations (UTO) model
- Sovereign-by-design architecture patterns
- Continuous Adversarial Validation (CAV) discipline

1.4 Target Audience

- Chief Information Security Officers (CISOs)
- Security operations centre (SOC) directors
- Security architects and platform engineers
- Detection engineering teams

© Intelligent Consulting 2026

- Regulatory compliance officers
- Board-level risk committees

© Intelligent Consulting 2026

2. Architectural Overview

2.1 Six-Layer Architecture

AURIS replaces the fragmented SIEM/SOAR/EDR/NDR/TIP stack with six integrated operational layers[1]:

Layer	Function	Impact
L1: STRATIX Mesh	Cloud-native telemetry data layer	60% cost ↓
L2: Intelligence	ML + AI detection	70% FP ↓
L3: Autonomy	Agentic operations	80% resolution
L4: UTO	Unified threat operations	IT/OT/Cloud/AI
L5: Sovereign	EU compliance architecture	NIS2/DORA auto
L6: CAV	Continuous validation	80% coverage

Table 1: AURIS Six-Layer Architecture

2.2 Design Philosophy

Intelligence-First: Detection embedded in schema, not downstream enrichment[1].

Autonomous by Default: 80% of alerts handled end-to-end without human intervention[1].

Sovereign by Design: Architecturally sovereign, not sovereign-by-configuration[1].

Continuously Validated: Detection coverage tested against live adversarial simulation[1].

© Intelligent Consulting 2026

3. Layer 1: STRATIX Telemetry Mesh

3.1 Overview

The STRATIX Telemetry Mesh is the foundational data layer that ingests, normalises, enriches, and routes all security telemetry in real time, with STRATIX as the universal schema across every data source and domain[1].

3.2 Architecture Components

Cloud-Native Data Layer:

- Object storage (S3-compatible, Azure Data Lake, Databricks Delta Lake)
- Open table formats (Apache Iceberg, Delta Lake, Apache Hudi)
- Separation of storage and compute
- Elastic scaling for query and ingestion workloads

STRATIX Normalisation Pipeline:

1. Raw telemetry ingestion from all sources
2. Vendor-specific parsing (ECS, CIM, ASIM, proprietary formats)
3. STRATIX normalisation at pipeline boundary
4. Intent layer enrichment (MITRE ATT&CK context)
5. Sovereignty metadata tagging (NIS2/DORA/GDPR)
6. Schema validation and routing

3.3 Telemetry Sources

Endpoint Telemetry:

- Process execution (create, inject, terminate)
- File system operations (write, delete, modify)
- Registry modifications (Windows)
- Network connections (outbound, lateral)
- Authentication events (local, remote)

Network Telemetry:

- DNS queries and responses
- Proxy logs (web traffic, blocked categories)
- NetFlow/IPFIX (traffic patterns, volume analysis)
- TLS certificate metadata

© Intelligent Consulting 2026

- Firewall allow/deny decisions

Identity Telemetry:

- Authentication attempts (success, failure, MFA)
- Privilege escalation events
- Service account behaviour
- Access token generation
- Password changes and resets

Cloud Telemetry:

- API calls (CloudTrail, Azure Activity Log, GCP Audit Logs)
- IAM policy modifications
- Configuration changes (security groups, storage permissions)
- Container and serverless execution logs
- Cloud storage access patterns

OT/ICS Telemetry:

- PLC state changes (STRATIX Layer 2)
- SCADA alarms and acknowledgements
- Industrial protocol events (Modbus, DNP3, OPC-UA)
- Engineering workstation actions
- Zone crossing violations (Purdue Model)

AI Agent Telemetry:

- Model invocations (STRATIX Layer 4)
- Tool usage and autonomous actions
- Decision traces and reasoning chains
- Human escalations
- Model drift signals

3.4 Cost Optimisation

Pipeline intelligence reduces ingestion costs by 40–60%[1]:

- **Pre-ingestion filtering:** Eliminate noise before storage (verbose debugging logs, known-good processes)
- **In-flight enrichment:** Add threat intelligence context during ingestion, not at query time
- **Classification-based routing:** Route high-value signals to hot storage, low-priority to cold tiers using STRATIX sovereignty.classification

© Intelligent Consulting 2026

- **Compression and deduplication:** Schema-aware compression (60–75% storage reduction)

3.5 Data Ownership and Control

Absolute data sovereignty:

- Your telemetry under your encryption keys
- Storage within your chosen jurisdiction
- No vendor lock-in (open table formats enable migration)
- Independent compute scaling (cost optimisation without vendor constraints)

© Intelligent Consulting 2026

4. Layer 2: Intelligence-First Detection

4.1 Overview

Detection in AURIS is not a ruleset. It is a continuously learning, multi-method intelligence layer operating across three complementary approaches simultaneously[1].

4.2 Behavioural Baseline Analysis

Baseline normal behaviour for every entity:

- Users (authentication patterns, geo-velocity, access times)
- Devices (process execution baselines, network connections)
- Service accounts (API call patterns, credential usage)
- Workloads (resource consumption, inter-service communication)

Statistically anomalous deviations surfaced against STRATIX intent layer:

Detection occurs not because a rule triggered, but because the system understands what normal looks like for this specific entity in this operational context[1].

Example: User account authenticates from new country + accesses sensitive file shares + during non-business hours = high confidence intent.category: "Initial Access" with intent.technique_id: "T1078" (Valid Accounts).

4.3 Continuous Machine Learning Detection

Pre-built and custom ML models operating continuously over STRATIX Telemetry Mesh:

- Authentication anomaly detection (impossible travel, credential stuffing)
- Lateral movement pattern recognition (RDP/SMB chaining)
- Data exfiltration volume analysis (DNS tunnelling, large transfers)
- Command and control beaconing detection (periodic network connections)
- Privilege escalation sequence detection (token manipulation patterns)

Schema normalisation advantage:

A single ML model trained on authentication patterns operates identically whether the source is Active Directory, Entra ID, Okta, or an OT HMI authentication log, because STRATIX normalises all to consistent schema[1].

STRATIX intent.confidence_score provides standardised, combinable signal across all detection methods[1].

© Intelligent Consulting 2026

4.4 AI-Augmented Chain Correlation

Large language models operating over STRATIX-normalised telemetry:

- Identify attack chains spanning multiple signals, data sources, timeframes
- Reconstruct kill chains from STRATIX `intent.kill_chain_phase` sequences
- Generate natural language attack narratives for analyst consumption
- Reason over structured intelligence (STRATIX fields) rather than raw log text

Result: 70% reduction in false positives[1].

Analysts direct expertise at genuine threats, not noise. Investigation time shifts from "is this real?" to "what's the business impact?"

4.5 Threat Intelligence Operationalisation

Continuous evaluation of threat intelligence feeds against live STRATIX mesh:

- STIX/TAXII feed ingestion (indicators of compromise)
- Matching against `intent.technique_id` fields in telemetry
- Automatic detection content generation when new IOC matches existing technique
- Intelligence receipt to deployed detection: <24 hours[1]

Traditional SIEM: Analyst manually writes rule, tests, deploys (weeks).

AURIS: Intelligence feed triggers automated rule generation, validation in CAV layer, production deployment (hours).

© Intelligent Consulting 2026

5. Layer 3: Autonomous Operations

5.1 Overview

AURIS operates through autonomous, agentic intelligence — not AI that summarises alerts or suggests queries, but AI agents that plan, investigate, and respond with bounded, auditable autonomy[1].

This is where AURIS diverges most radically from legacy security market offerings.

5.2 Multi-Agent Orchestration Architecture

Specialised agents work in parallel and coordination:

Agent	Function
Triage Agent	Receives all STRATIX alerts, scores against business context (asset criticality from sovereignty layer, NIS2/DORA classification), autonomously routes to escalation, auto-resolution, or deeper investigation
Investigation Agent	Gathers evidence by querying STRATIX Telemetry Mesh, pulls threat intelligence against intent.technique_id fields, reconstructs kill chain from intent.kill_chain_phase sequences, produces case file with verdict and recommended action
Response Agent	Executes bounded containment actions within pre-approved governance guardrails (isolate endpoint, disable compromised account, block malicious IP, create firewall rule). Every action recorded as ai.autonomous_action event in STRATIX
Proactive Hunt Agent	Operates continuously, runs hypothesis-driven investigations across full STRATIX mesh, leverages structured intent layer to surface weak-signal attack chain fragments
Orchestration Agent	Coordinates all agents, manages workflows, escalates to human analysts when uncertainty exceeds defined thresholds. All escalations recorded as ai.human_escalation events

Table 2: AURIS Agent Framework

5.3 Governance Framework

Autonomy operates within defined boundaries:

- **Fully autonomous actions:** Alert triage, evidence gathering, low-risk containment (non-production asset isolation)

© Intelligent Consulting 2026

- **Human-approved actions:** Production asset isolation, account disablement, firewall rule creation
- **Human-only decisions:** Legal/regulatory escalation, board notification, law enforcement engagement

All autonomous actions captured in STRATIX Layer 4 (AI Telemetry):

Every decision, tool call, and reasoning chain preserved as immutable audit trail, providing the governance record required by EU AI Act transparency and accountability requirements[1].

5.4 The Aviation Autopilot Model

Commercial aviation provides the operational model[1]:

- Autopilot operates the aircraft 95% of the time with superhuman precision
- Pilots (highly trained, fully accountable) intervene when conditions exceed automation envelope
- Clear governance boundaries define when human judgement is required

AURIS is built on exactly this principle.

80% of alerts handled autonomously end-to-end. Human analysts operate exclusively at investigation and strategic levels, not manual triage[1].

5.5 Performance Metrics

Target outcomes (validated across multiple deployments):

Metric	Legacy SIEM	AURIS Target
Autonomous resolution rate	0%	80%[1]
Mean time to triage	Hours	<5 minutes
Mean time to investigate	Days	<30 minutes
Mean time to respond	Hours	Seconds (autonomous)
Analyst focus	L1 triage	L3 hunt/strategy

Table 3: Autonomous Operations Performance

© Intelligent Consulting 2026

6. Layer 4: Unified Threat Operations (UTO)

6.1 Overview

Unified Threat Operations (UTO) is an original AURIS framework concept — the cross-domain correlation and response discipline that replaces the fragmented SIEM/SOAR/EDR/NDR/CSPM/CIEM/TIP stack with a single, coherent operational layer built on the STRATIX Telemetry Mesh[1].

UTO is not a vendor product category. It is an operational model.

6.2 Cross-Domain Correlation

All telemetry domains operate as unified fabric:

Endpoint ↔ Network ↔ Identity ↔ Cloud:

An attack that begins at the endpoint (process injection), pivots through a compromised identity (token theft), and exfiltrates via a cloud storage misconfiguration surfaces as a single, coherent STRATIX kill chain — not three separate alerts in three separate consoles[1].

Example attack chain:

1. Endpoint: Process injection detected (`intent.technique_id: "T1055"`)
2. Identity: Token theft (`intent.technique_id: "T1528"`)
3. Network: C2 beaconing to known malicious IP (`intent.technique_id: "T1071.001"`)
4. Cloud: Abnormal S3 data exfiltration (`intent.technique_id: "T1567.002"`)

UTO correlation: Single incident timeline, unified kill chain reconstruction, automated response across all four domains.

6.3 OT/ICS as First-Class Participant

STRATIX Layer 2 event classes mean OT telemetry correlates natively with IT/Cloud[1]:

Example OT attack chain:

1. Engineering workstation: Anomalous authentication (`ot.engineering_workstation_action`)
2. Network: DNS query to C2 domain from OT zone (`intent.technique_id: "T0884"`)

© Intelligent Consulting 2026

3. PLC: Unauthorised program upload (`ot.plc_state_change`, `intent.technique_id: "T0855"`)
4. Process: Parameter deviation triggering safety alarm (`ot.process_deviation`)

UTO surfaces this as unified cross-domain incident, not isolated IT and OT alerts.

For NIS2 Article 21 compliance (security of OT environments for operators of essential services), this unified visibility is mandatory[1].

6.4 Cloud Security Posture Management (CSPM) Native

CSPM is operational function within UTO, not separate tool:

- Misconfigured S3 buckets, Azure storage, GCS objects detected proactively
- Over-permissioned service accounts (CIEM) surfaced with IAM policy context
- Public-facing workloads identified and risk-scored
- Enriched with STRATIX sovereignty metadata to immediately flag GDPR or NIS2 exposure

No separate CSPM console. No duplicate telemetry ingestion. Unified operational layer.

6.5 UTO Operational Workflow

Single pane of glass for all threat operations:

1. Alert generation (any domain: IT, OT, Cloud, AI)
2. Autonomous triage (Triage Agent scores with business context)
3. Cross-domain correlation (UTO reconstructs full kill chain)
4. Automated investigation (Investigation Agent gathers evidence)
5. Bounded response (Response Agent executes containment)
6. Human escalation only when required (Orchestration Agent threshold)

Analyst experience: One incident timeline, one kill chain visualisation, one response workflow — regardless of attack surface span.

© Intelligent Consulting 2026

7. Layer 5: Sovereign-by-Design Architecture

7.1 Overview

AURIS is sovereign-by-design, not sovereign-by-configuration or sovereign-by-contractual-assurance, but architecturally sovereign from the ground up[1].

This is non-negotiable for European organisations operating in regulated sectors or handling sensitive data.

7.2 Sovereignty Dimensions

Dimension	What It Means	STRATIX Field	Regulatory Obligation
Data Residency	Where data is physically stored	sovereignty.data_residency	GDPR data transfer restrictions
Data Classification	Sensitivity and handling requirements	sovereignty.classification	GDPR Article 5, NIS2 Article 21
Jurisdictional Control	Which legal system governs operations	sovereignty.eucs_assurance_level	EUCS certification, US CLOUD Act exposure
Regulatory Scope	Which obligations apply to this event	sovereignty.nis2_category, sovereignty.dora_ict_asset	NIS2 notification timelines, DORA resilience testing
AI Governance	AI system classification	sovereignty.ai_act_classification	EU AI Act transparency and accountability

Table 4: Sovereignty Architecture Dimensions

7.3 US CLOUD Act Risk

The most underappreciated risk in European security architecture[1]:

A US-headquartered cloud provider can be legally compelled to provide access to EU-hosted data, regardless of contractual data residency commitments.

© Intelligent Consulting 2026

Every organisation routing security telemetry (the most sensitive operational data it generates) through a US-controlled cloud platform carries this exposure.

STRATIX `sovereignty.access_log` provides immutable, cryptographically signed record of every access to every event record, giving organisations the evidence base to detect and challenge unlawful access attempts[1].

7.4 Sovereign Architecture Requirements

AURIS sovereign deployment requires:

- **EU-operated cloud infrastructure:** STRATIX Telemetry Mesh deployed on EUCS-certified infrastructure operated by EU legal entities, under EU jurisdiction, with no non-EU access paths
- **Sovereign cloud providers:** OVHcloud, Deutsche Telekom Open Telekom Cloud, Hetzner, Impossible Cloud European Sovereign Cloud offerings
- **Customer-controlled encryption:** Bring-your-own-key (BYOK) or hold-your-own-key (HYOK) models
- **AI inference sovereignty:** All autonomous agent inference within sovereign boundary. Locally deployed open-weight models (Mistral, LLaMA derivatives) operating within EUCS-certified boundary

7.5 NIS2/DORA/GDPR/AI Act Alignment by Default

STRATIX Layer 3 metadata makes compliance reporting automatable[1]:

NIS2 Article 23 notifications:

- 24-hour early warning: Auto-generated from `sovereignty.nis2_category: "Essential Entity"` incidents
- 72-hour full report: Pre-populated from STRATIX event records, kill chain reconstruction, impact assessment

DORA Article 19 ICT incident reporting:

- Major ICT incidents: Automatically flagged when `sovereignty.dora_ict_asset: true` affected
- Resilience testing records: CAV simulation results provide DORA testing evidence

GDPR Article 30 processing records:

- `sovereignty.access_log` provides complete audit trail of data processing activities
- `sovereignty.gdpr_lawful_basis` documents legal basis for every event

EU AI Act Articles 13 & 14:

© Intelligent Consulting 2026

- `ai.decision_trace` preserves reasoning chains for transparency
- `ai.autonomous_action` provides accountability records
- `ai.governance_audit_record` generates regulatory audit packages

Manual analyst hours for compliance reporting: 1,200+ → 0[1].

© Intelligent Consulting 2026

8. Layer 6: Continuous Adversarial Validation (CAV)

8.1 Overview

A platform never tested against real adversarial conditions is a platform that fails silently.

AURIS embeds Continuous Adversarial Validation (CAV) as native operational discipline — not a periodic exercise, not an annual penetration test, but continuously running validation layer integrated into normal operations[1].

8.2 CAV Operational Model

Automated adversarial simulation:

- Continuously executing MITRE ATT&CK techniques against live environment
- Validates detection coverage and response fidelity
- Gaps surface as STRATIX events with `intent.technique_id` populated
- Coverage gaps immediately queryable and actionable

Full kill chain simulation:

- Automated simulation of attacker behaviours from initial access through impact
- Spans complete MITRE ATT&CK kill chain
- Simulation events flagged in STRATIX to distinguish from genuine threats

8.3 Detection-as-Code

All detection logic managed in version-controlled repositories:

- Detection rules, ML model configurations, response workflows as code
- Tested in continuous integration/deployment pipelines before production release
- No unvalidated detection logic accumulating technical debt
- Every change peer-reviewed, tested against CAV simulation

Detection engineering velocity: Weeks → <24 hours[1].

8.4 Operationalised Threat Intelligence

Threat intelligence feeds evaluated continuously against live STRATIX mesh:

© Intelligent Consulting 2026

- When new indicator or technique matches `intent.technique_id` in event stream, detection content automatically generated or updated
- Intelligence receipt to detection deployment: <24 hours[1]
- Closes gap between "we know about this threat" and "we can detect this threat"

8.5 MITRE ATT&CK Coverage Metrics

Platform	Coverage	Validation Method
Legacy SIEM	15–25%	Annual penetration test
AURIS + CAV	80%[1]	Continuous simulation

Table 5: MITRE ATT&CK Coverage Comparison

Coverage becomes the primary, publicly reported measure of detection maturity[1].

© Intelligent Consulting 2026

9. Platform Implementation Models

9.1 AURIS Platform Agnosticism

AURIS is platform-agnostic by design. The STRATIX Telemetry Mesh can be deployed over any cloud-native security analytics infrastructure[1].

9.2 Implementation Approaches

Platform Approach	Representative Vendors	STRATIX Fit	Sovereignty Posture
Search + AI Security Analytics	Elastic Security	High (open architecture, ECS-to-STRATIX mapper, sovereign deployment flexibility)	Deployable on EU sovereign infrastructure (open source core)
Intelligence-First Security Analytics	CrowdStrike Falcon, SentinelOne Singularity	Medium (strong endpoint telemetry, STRATIX integration via API pipeline)	US-headquartered, sovereign deployment limited
Unified Security Platform	Palo Alto Cortex	Medium (strong platform breadth, STRATIX ingestion via pipeline adapters)	US-headquartered, EU region available
Cloud-Native Security Analytics	Microsoft Sentinel	Medium (deep M365 integration, STRATIX via ASIM mapper)	EU Data Boundary complete 2025, CLOUD Act exposure remains
EU-Native SOC Platform	SEKOIA.IO, Gatewatcher	High (EU-headquartered, sovereign-first, STRATIX alignment natural)	EU-headquartered and operated
Federated Security Analytics	Query.AI	High (decentralised architecture aligns with STRATIX sovereignty model)	Architecture-level sovereignty support

Table 6: AURIS Implementation Platform Options

© Intelligent Consulting 2026

9.3 Vendors to Challenge

The incumbents to challenge hardest:

- IBM QRadar
- Splunk Enterprise Security
- ArcSight (Micro Focus/OpenText)

Why: AI additions are bolt-on to volume-priced, proprietary architectures. None can natively consume STRATIX metadata — integration requires custom development that recreates the very debt organisations are seeking to escape[1].

© Intelligent Consulting 2026

10. The Human Dimension: Redefining the Analyst Role

10.1 AURIS Does Not Eliminate Analysts

AURIS elevates them[1].

Level 1 analyst roles absorbed into AURIS Autonomous Operations Layer:

- Manual alert triage
- Basic enrichment
- Ticket creation
- First-response checklist execution

10.2 The AURIS SOC Operating Model

Analysts retrained and repositioned:

Detection Engineers:

- Building, testing, validating STRATIX-native detection logic as code
- Owning MITRE ATT&CK coverage strategy
- Curating intent layer vocabulary for organisation's threat profile

Threat Intelligence Analysts:

- Operationalising threat intelligence against STRATIX `intent.technique_id` fields
- Producing finished intelligence for autonomous detection systems
- Engaging with sector-specific information sharing communities

Proactive Hunt Analysts:

- Hypothesis-driven, continuous investigation across full STRATIX Telemetry Mesh
- The highest-value analytical function in the AURIS SOC

AI Governance Analysts (genuinely new role):

- Owning autonomy governance framework
- Auditing `ai.decision_trace` and `ai.autonomous_action` records
- Managing human-in-the-loop thresholds
- Producing EU AI Act compliance audit reports

STRATIX Schema Architects:

- Owning organisation's STRATIX implementation profile
- Managing domain-specific extensions

© Intelligent Consulting 2026

- Ensuring conformance posture
- Evolving schema as threat and regulatory landscape changes

The SOC of 2028 is not smaller. It is more strategic, more capable, and operating at qualitatively higher level of analytical sophistication than anything the legacy SIEM era could support[1].

© Intelligent Consulting 2026

11. Economics and Business Case

11.1 Total Cost of Ownership Analysis

Cost Component	Legacy SIEM	AURIS STRATIX	Reduction
Infrastructure	€250,000	€50,000	80%
Ingestion/Licensing	€300,000	€120,000	60%
Detection Engineering	€150,000	€60,000	60%
Analyst Time (Triage)	€240,000	€60,000	75%
Compliance Reporting	€120,000	€0	100%
Total Annual Cost	**€1,060,000**	**€290,000**	**73%**

Table 7: TCO Comparison (5,000-person organisation)

Target TCO reduction: 60–75% below legacy stack baseline[1].

11.2 Business Case by Stakeholder

For the CISO and Security Architect:

Legacy SIEM in-house: €380k–€940k per year. Full fragmented stack at enterprise scale: €5.6M–€9.4M per year[1].

AURIS on STRATIX Telemetry Mesh: 60–75% lower total cost of ownership, with dramatically higher detection fidelity, reduced mean time to respond, native regulatory compliance (no custom development required)[1].

For the CFO and Board:

A single material breach reportable under NIS2 (€10M or 2% of global turnover), DORA (potential operational suspension for financial entities), or GDPR (€20M or 4% of global turnover) dwarfs the total cost of platform modernisation[1].

STRATIX automated compliance reporting eliminates manual evidence-gathering burden that currently costs organisations average of 1,200 analyst-hours per reportable incident[1].

The legacy SIEM is not a risk mitigation investment — it is a risk accumulation one[1].

For the CEO and Risk Committee:

Sovereignty risk is now a business risk, not just a compliance risk[1].

Security telemetry flowing through platform subject to non-EU jurisdictional control creates permanent, unhedgeable legal exposure.

© Intelligent Consulting 2026

In geopolitical environment where trans-Atlantic data frameworks face existential legal challenge, this is strategic liability of board-level consequence[1].

STRATIX sovereignty.access_log is the only schema-native mechanism that provides cryptographic proof of data access control — the foundation of legally defensible, audit-ready sovereignty posture[1].

© Intelligent Consulting 2026

12. Implementation Roadmap

12.1 Four-Phase Transformation Programme

This is not a rip-and-replace exercise. It is structured operational transformation programme[1].

Phase 1: Foundation (Months 1–6)

Objectives:

- Build STRATIX Telemetry Mesh
- Deploy security data layer on sovereign-compliant cloud infrastructure
- Implement STRATIX normalisation at pipeline boundary, beginning with sovereignty metadata layer
- Migrate high-fidelity, high-volume sources first (endpoint, identity, cloud)
- Run legacy SIEM in parallel (do not decommission until STRATIX coverage independently validated)
- Establish Detection-as-Code practices (version control all detection logic from day one)

Cost Impact:

Pipeline optimisation to STRATIX schema typically reduces ingestion costs by 40–60% within this phase[1].

Phase 2: Intelligence (Months 6–12)

Objectives:

- Activate Intelligence-First Detection
- Deploy continuous machine learning detection across all STRATIX-normalised telemetry
- Implement behavioural baseline analysis leveraging STRATIX intent and entity context
- Integrate threat intelligence operationalisation against `intent.technique_id` fields
- Deploy AI-augmented chain correlation operating over STRATIX kill chain layer
- Establish MITRE ATT&CK coverage as primary, publicly reported measure of detection maturity
- Add STRATIX Layer 2 OT event classes for all OT/ICS environments in scope

Outcome Targets:

- 70% reduction in false positive volume[1]
- 60% MTTD reduction

© Intelligent Consulting 2026

Phase 3: Autonomy (Months 12–18)

Objectives:

- Deploy AURIS Autonomous Operations Layer
- Define autonomy governance framework (which decisions are fully autonomous, which require human approval, which remain human-only)
- Deploy Triage Agent (eliminate Level 1 manual alert triage as human function)
- Implement Response Agent for approved containment actions (every action captured as `ai.autonomous_action` in STRATIX)
- Preserve all reasoning as `ai.decision_trace` records (EU AI Act audit-ready from day one)
- Activate Continuous Adversarial Validation (CAV) across live environment

Outcome Targets:

- 80% of alerts handled autonomously within governance policy[1]
- Human analysts operating exclusively at investigation and strategic levels

Phase 4: Optimisation (Months 18–24)

Objectives:

- Complete legacy SIEM decommission (eliminate all associated infrastructure, licensing, operational overhead)
- Achieve EUCS High alignment (validated through STRATIX `sovereignty.eucs_assurance_level` audit records)
- Implement full DORA and NIS2 compliance reporting automation from STRATIX metadata
- Conduct first full sovereign architecture audit (validate zero data flows outside STRATIX `sovereignty.data_residency` boundaries)
- Engage Intelligent Consulting BV for AURIS implementation advisory and STRATIX conformance assessment

Outcome Target:

Full sovereign compliance posture. TCO 60–75% below legacy stack baseline[1].

© Intelligent Consulting 2026

13. Performance Metrics and SLAs

13.1 The New Metrics of Security Operations Excellence

Metric	Legacy SIEM Baseline	AURIS STRATIX Target
Mean Time to Detect (MTTD)	Hours to days	<5 minutes (high-severity)[1]
Mean Time to Respond (MTTR)	Hours	Seconds (autonomous), minutes (full remediation)[1]
False Positive Rate	40–60% of all alerts	<5%[1]
MITRE ATT&CK Coverage	15–25% validated	80% continuously validated via CAV[1]
Autonomous Resolution Rate	0%	80% of alerts end-to-end[1]
NIS2 Reporting Time	Days of manual effort	Auto-generated from STRATIX sovereignty metadata[1]
Sovereignty Audit Readiness	Manual (months of preparation)	Always-on (STRATIX access logs are audit record)[1]
Detection Engineering Velocity	Weeks per rule	<24 hours (new intelligence to deployed detection)[1]

Table 8: AURIS Performance Metrics

© Intelligent Consulting 2026

14. Security and Resilience

14.1 Platform Security Architecture

AURIS implements defence-in-depth:

- Field-level encryption for sensitive STRATIX fields
- Role-based access control (RBAC) integrated with `sovereignty.access_log`
- Immutable audit trails (append-only logs, cryptographically signed)
- Zero-trust network architecture (all inter-service communication authenticated)
- AI agent sandboxing (bounded execution environments for autonomous actions)

14.2 Resilience and Business Continuity

High availability architecture:

- Multi-region deployment for sovereign resilience
- Automated failover (RPO <5 minutes, RTO <15 minutes)
- Continuous backup of STRATIX Telemetry Mesh
- Disaster recovery testing as part of CAV discipline

© Intelligent Consulting 2026

15. Standards and Certifications

15.1 Compliance Framework Support

Standard/Regulation	AURIS Layer	Compliance Approach
NIS2 (EU 2022/2555)	L5 Sovereign	Automated 24h/72h reporting
DORA (EU 2022/2554)	L5 Sovereign	ICT resilience automation
GDPR (EU 2016/679)	L5 Sovereign	Article 30 audit trails
EU AI Act (EU 2024/1689)	L4 AI Telemetry	Transparency & accountability
IEC 62443 (OT Security)	L2 STRATIX	Native OT event classes
MITRE ATT&CK	L1 STRATIX	Intent layer embedding
ISO 27001	All layers	Controls mapping
SOC 2 Type II	All layers	Audit evidence generation

Table 9: Standards and Regulatory Compliance

15.2 Certification Roadmap

Target certifications for AURIS platform:

- EUCS High (EU Cybersecurity Certification Scheme)
- Common Criteria EAL4+ (Security Operations Platform)
- ISO 27001 (Information Security Management)
- SOC 2 Type II (Security, Availability, Confidentiality)

© Intelligent Consulting 2026

16. Vendor Ecosystem and Partnerships

16.1 Strategic Technology Partners

STRATIX schema adoption partners:

- Elastic (ECS-to-STRATIX mapper)
- Microsoft (ASIM-to-STRATIX integration)
- Splunk (CIM-to-STRATIX adapter)
- OT vendors (Claroty, Nozomi, Dragos protocol adapters)

Sovereign cloud infrastructure partners:

- OVHcloud (French sovereign hosting)
- Deutsche Telekom Open Telekom Cloud (German sovereign)
- Hetzner (German data centre operations)
- Impossible Cloud (European distributed sovereign cloud)

AI inference partners:

- Mistral AI (European LLM provider)
- Aleph Alpha (German sovereign AI)
- Open-weight model deployment (LLaMA, Falcon within sovereign boundary)

© Intelligent Consulting 2026

17. Conclusion: The Strategic Imperative

SIEM, as a concept, is not dead, but as an implementation architecture rooted in centralised log aggregation, rule-based correlation, proprietary data formats, and volume-based pricing, it is obsolete[1].

Defending that architecture is not conservatism — it is negligence.

AURIS, built on STRATIX Telemetry Mesh, powered by intelligence-first detection and Autonomous Operations Layer, unified across full operational surface (IT/OT/Cloud/AI agents) through Unified Threat Operations, continuously validated through Continuous Adversarial Validation, and sovereign-by-design through STRATIX native regulatory metadata, is not a vision for 2030[1].

The components exist today. The framework is available now.

The regulatory imperatives in Europe make sovereign architecture mandatory, not optional.

Organisations that begin this transition now will have mature, autonomous, sovereign-compliant security operations capability by 2028.

Those that continue to invest in legacy SIEM modernisation will be fighting the same battles, paying the same bills, absorbing the same breach risk, with an ever-wider gap of missed architectural generation to close[1].

The question is not whether to make this transition. The question is whether you lead it or are forced into it by a breach that could have been prevented[1].

© Intelligent Consulting 2026

References

[1] Button, S. N. (2026). *SIEM Is Dead! Long Live STRATIX and AURIS: A Strategic Vision for the Future of Security Operations*. Intelligent Consulting BV.

[2] Open Cybersecurity Schema Framework (OCSF). (2024). *OCSF Schema Version 1.1*. https://schema.ocsf.io/

[3] MITRE Corporation. (2024). *MITRE ATT&CK Framework for Enterprise*. https://attack.mitre.org/

[4] MITRE Corporation. (2024). *MITRE ATT&CK for ICS*. https://attack.mitre.org/matrices/ics/

[5] European Parliament and Council. (2022). *Directive (EU) 2022/2555 on measures for a high common level of cybersecurity across the Union (NIS2)*. Official Journal of the European Union.

[6] European Parliament and Council. (2022). *Regulation (EU) 2022/2554 on digital operational resilience for the financial sector (DORA)*. Official Journal of the European Union.

[7] European Parliament and Council. (2016). *Regulation (EU) 2016/679 on the protection of natural persons with regard to the processing of personal data (GDPR)*. Official Journal of the European Union.

[8] European Parliament and Council. (2024). *Regulation (EU) 2024/1689 laying down harmonised rules on artificial intelligence (EU AI Act)*. Official Journal of the European Union.

[9] International Electrotechnical Commission. (2018). *IEC 62443 Series: Security for industrial automation and control systems*. IEC Standards.

[10] European Union Agency for Cybersecurity (ENISA). (2024). *EU Cybersecurity Certification Scheme (EUCS)*. https://www.enisa.europa.eu/

[11] National Institute of Standards and Technology. (2018). *Framework for Improving Critical Infrastructure Cybersecurity Version 1.1*. NIST.

[12] Cloud Security Alliance. (2023). *Security Guidance for Critical Areas of Focus in Cloud Computing v5.0*. CSA.

© Intelligent Consulting 2026

Appendix A: AURIS Agent Governance Template

A.1 Autonomy Decision Matrix

Action Type	Risk Level	Approval	Rollback Time
Alert triage	Low	Autonomous	N/A
Evidence gathering	Low	Autonomous	N/A
Ticket creation	Low	Autonomous	Immediate
Non-production asset isolation	Medium	Autonomous	<5 minutes
Production asset isolation	High	Human approval	<15 minutes
Account disablement	High	Human approval	<5 minutes
Firewall rule creation	High	Human approval	<2 minutes
Legal/regulatory escalation	Critical	Human only	N/A

Table 10: Autonomy Governance Decision Matrix

A.2 Human-in-the-Loop Thresholds

Orchestration Agent escalates to human when:

- Confidence score <70 for high-impact action
- Multiple conflicting signals in kill chain reconstruction
- Asset criticality classification: "Critical Infrastructure" or "Sovereign"
- Legal/regulatory notification potentially required
- Action would impact >100 users or >10 critical systems

© Intelligent Consulting 2026

Appendix B: AURIS Deployment Checklist

B.1 Pre-Deployment Requirements

- [ ] STRATIX Telemetry Mesh infrastructure provisioned
- [ ] Sovereign cloud provider selected and contracted
- [ ] EUCS-certified hosting validated
- [ ] Customer-controlled encryption keys generated (BYOK/HYOK)
- [ ] Data residency boundaries documented
- [ ] Legacy SIEM parallel operation plan approved
- [ ] Detection-as-Code repository established
- [ ] Agent governance framework approved by legal/compliance

B.2 Phase 1 Foundation Validation

- [ ] STRATIX normalisation pipeline operational
- [ ] Sovereignty metadata enrichment validated
- [ ] High-fidelity sources migrated (endpoint, identity, cloud)
- [ ] Schema validation passing >99.5%
- [ ] Ingestion cost reduction measured (target: 40–60%)
- [ ] Legacy SIEM comparison coverage validated

B.3 Phase 2 Intelligence Activation

- [ ] Behavioural baseline models deployed
- [ ] ML detection jobs operational
- [ ] AI chain correlation active
- [ ] Threat intelligence feeds integrated
- [ ] False positive reduction measured (target: 70%)
- [ ] MITRE ATT&CK coverage baseline established

B.4 Phase 3 Autonomy Deployment

- [ ] Triage Agent operational
- [ ] Investigation Agent deployed

© Intelligent Consulting 2026

- [ ] Response Agent governance approved
- [ ] Autonomous action audit trails validated
- [ ] Human escalation thresholds tuned
- [ ] Autonomous resolution rate measured (target: 80%)

B.5 Phase 4 Optimisation Completion

- [ ] Legacy SIEM decommissioned
- [ ] EUCS High certification achieved
- [ ] NIS2/DORA reporting automation validated
- [ ] Sovereign architecture audit passed
- [ ] TCO reduction validated (target: 60–75%)
- [ ] MITRE ATT&CK coverage validated (target: 80%)

© Intelligent Consulting 2026

Appendix C: Glossary

AURIS — Autonomous Unified Response Intelligence System

CAV — Continuous Adversarial Validation

CSPM — Cloud Security Posture Management

CIEM — Cloud Infrastructure Entitlement Management

DORA — Digital Operational Resilience Act (EU 2022/2554)

EUCS — EU Cybersecurity Certification Scheme

MTTD — Mean Time to Detect

MTTR — Mean Time to Respond

NIS2 — Network and Information Security Directive 2 (EU 2022/2555)

OCSF — Open Cybersecurity Schema Framework

OT/ICS — Operational Technology / Industrial Control Systems

SIEM — Security Information and Event Management

SOAR — Security Orchestration, Automation and Response

STRATIX — Security Telemetry, Risk Attribution & Threat Intelligence eXchange

UTO — Unified Threat Operations

Document Control

Version	Date	Author	Changes
1.0	March 2026	S. N. Button	Initial publication

Copyright © 2026 Intelligent Consulting BV. All rights reserved.

This specification is protected under EU Directive 2019/790/EU and the Berne Convention. Reproduction or commercial use requires written authorisation from Intelligent Consulting BV.

For AURIS implementation advisory, STRATIX conformance assessment, or licensing enquiries:

Intelligent Consulting BV
Brussels, Belgium
Director: Suzanne Natalie Button
Email: contact@intelligentconsulting.eu